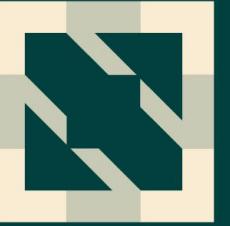


KubeCon



CloudNativeCon

S OPEN SOURCE SUMMIT

China 2023



KubeCon



CloudNativeCon



OPEN SOURCE SUMMIT

China 2023

# The Service Mesh Exploration of China Mobile Cloud

*Haiwen Zhang  
China Mobile*



KubeCon



CloudNativeCon



OPEN SOURCE SUMMIT

China 2023

# CONTENTS

**01**

**Application Practice**

**02**

**Technical Exploration**

**03**

**Product Development**

**04**

**Future Prospects**



KubeCon



CloudNativeCon



OPEN SOURCE SUMMIT

China 2023

# Application Practice

# Overview

## China Mobile Cloud

- China Mobile's cloud computing platform, TOP5 in public cloud market and TOP3 in the IaaS subdomain market in China.

## Scenarios

- Nearly 50 products are using service mesh for canary releases, traffic management, and observability.
- By OPS and CNP, we provide service mesh capabilities to both internal products and external customers.

### Applications

云专线

5G云梯

云视讯

...

### Platforms

OPS

CNP

### Scenarios

Canary Release

Governance

Observability

### Components

Prometheus

Kiali

Jaeger

Istio

### Infrastructures

BC-EKI

BC-CIS

Docker

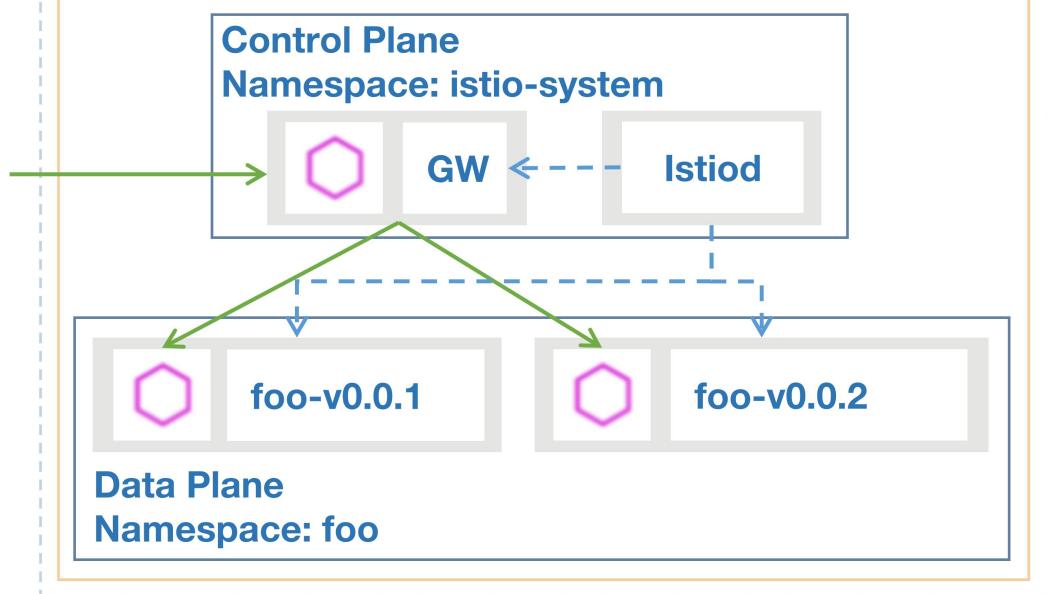
Containerd

# Canary Release

Many scenarios for canary release

## Istio Mesh

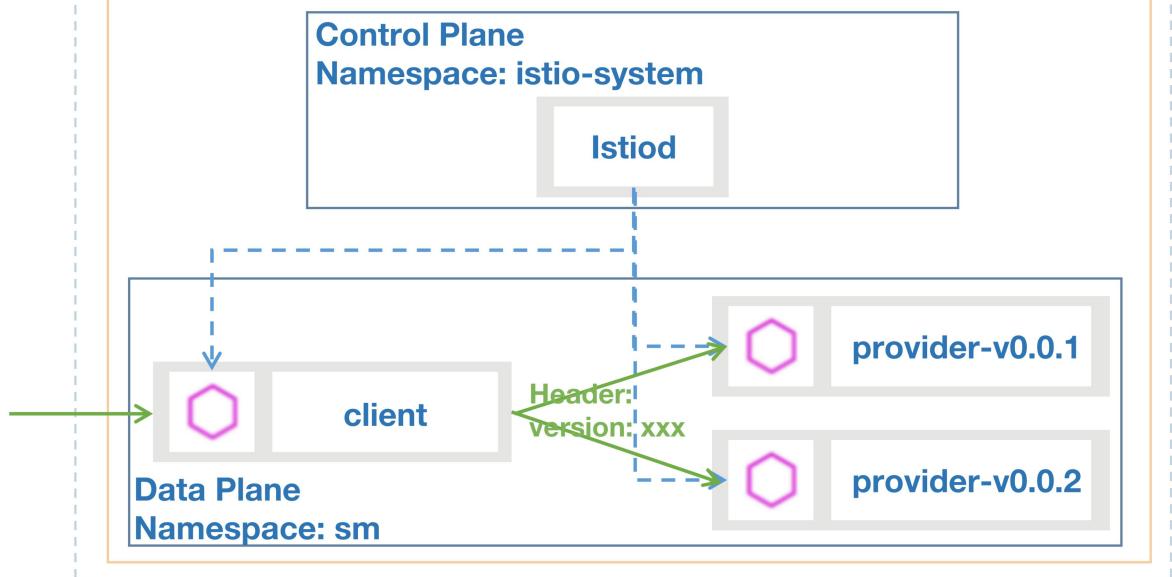
### Cluster



Canary Release for the Entry Microservice

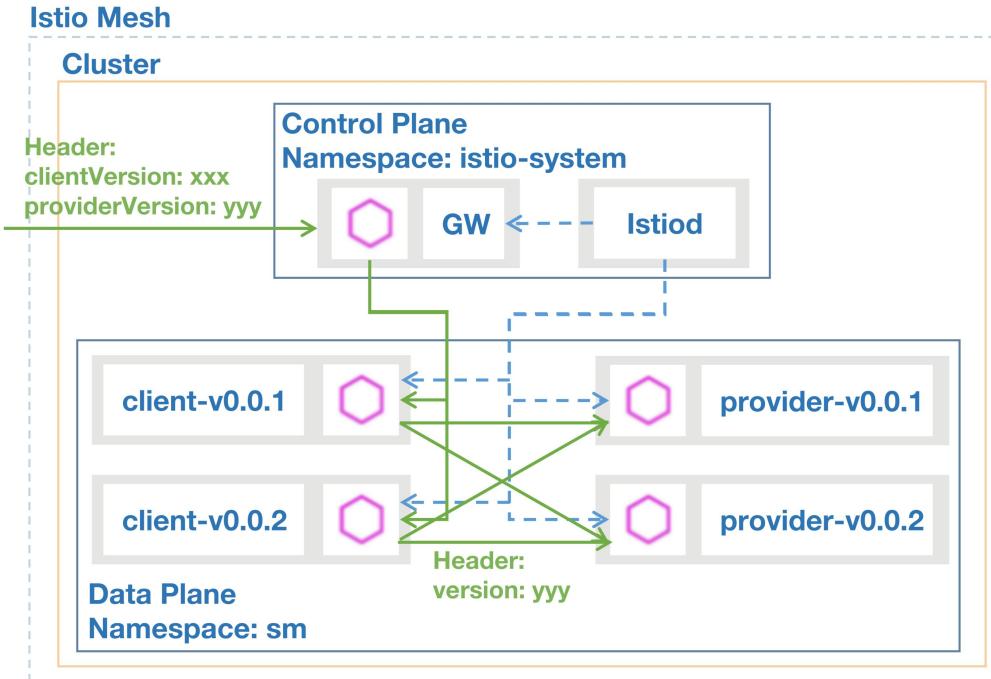
## Istio Mesh

### Cluster

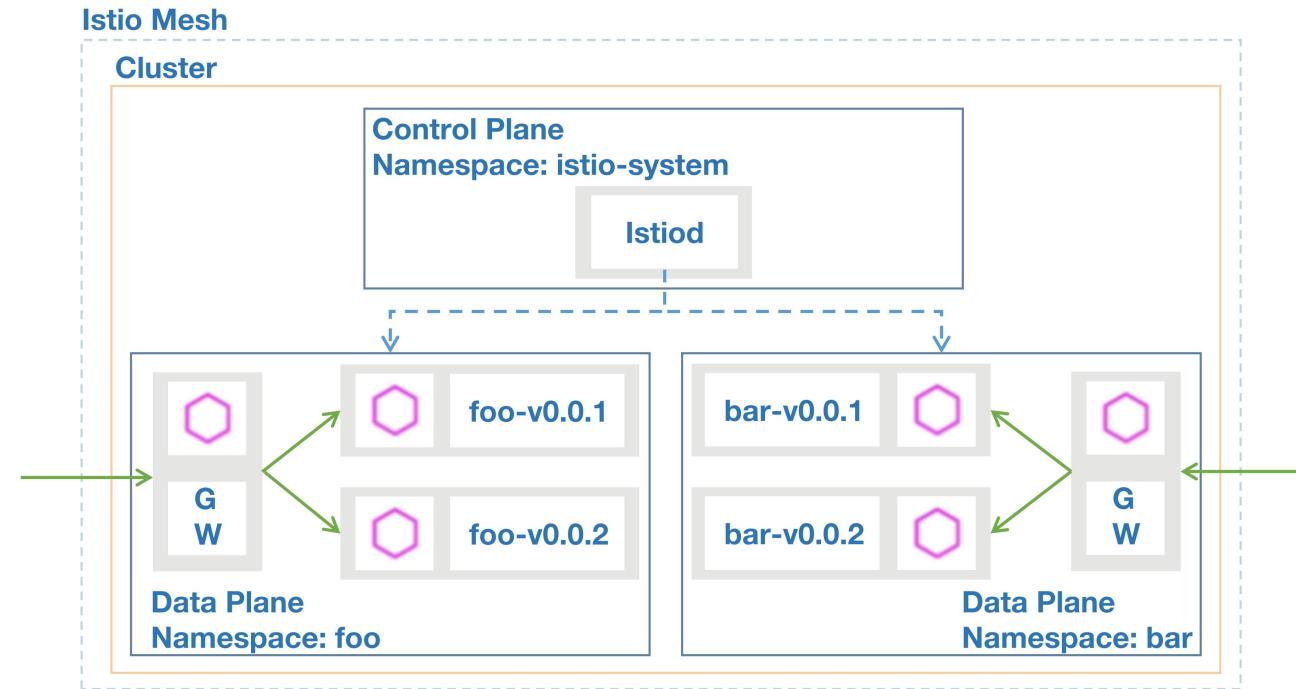


Canary Release for Internal Microservices

# Canary Release



Canary Release for End-to-End Microservices

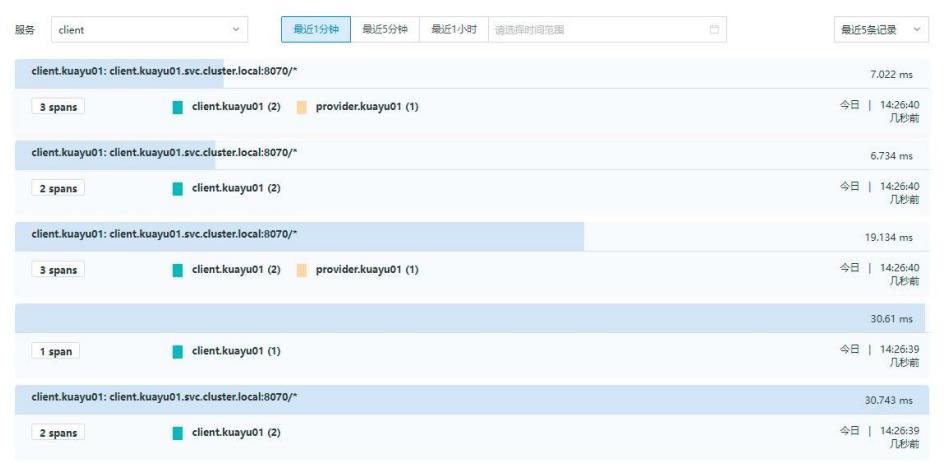


Canary Release for Multi-Tenant Microservices

# Observability



## Service Topology

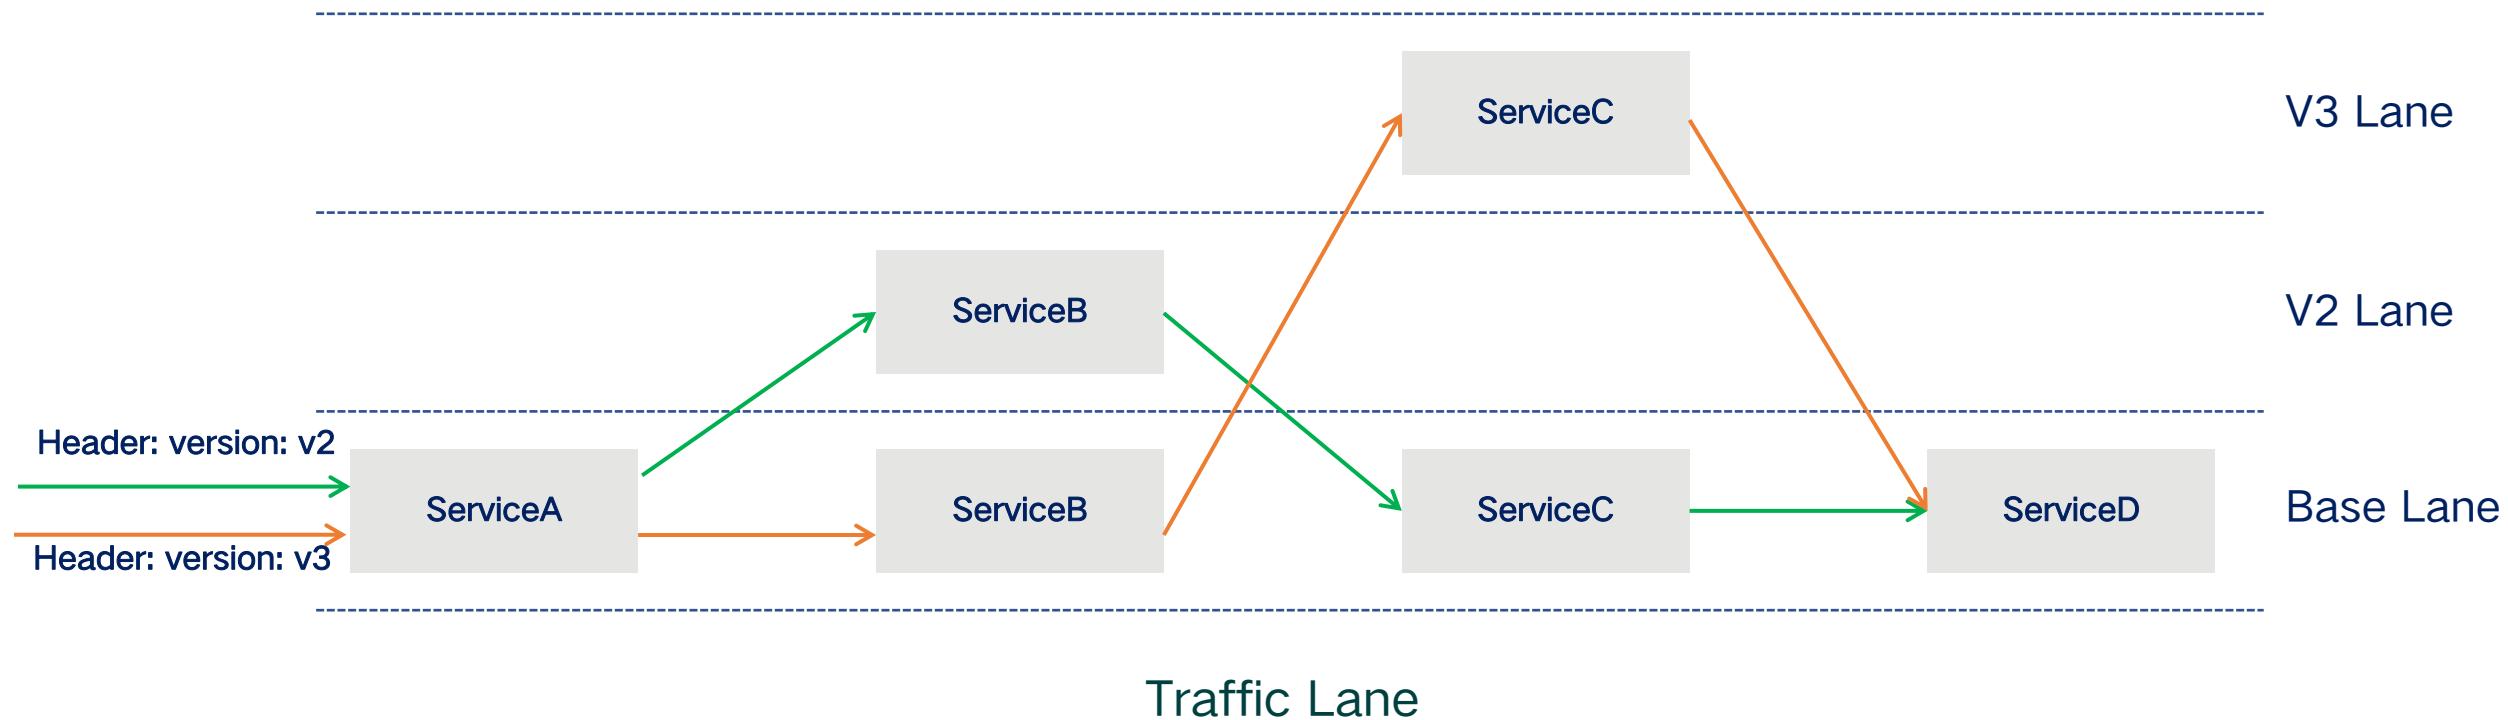


## Tracing

# Traffic Lane

**Background:** Dozens of microservices; Multiple concurrent development iterations; Challenge to Deploy the environment for each iteration; Sharing the same environment leads to mutual interference.

**Solution:** Building end-to-end traffic lanes based on Istio+OpenTelemetry to achieve iteration environment isolation.



# Computing Network-Native



## East-West Computing Project

Aim to enable low-cost computing resources in the western regions to support data processing in the eastern regions.

## Computing Network

A deep integration of computation and networking, to meet the demands of East-West Computing Project.

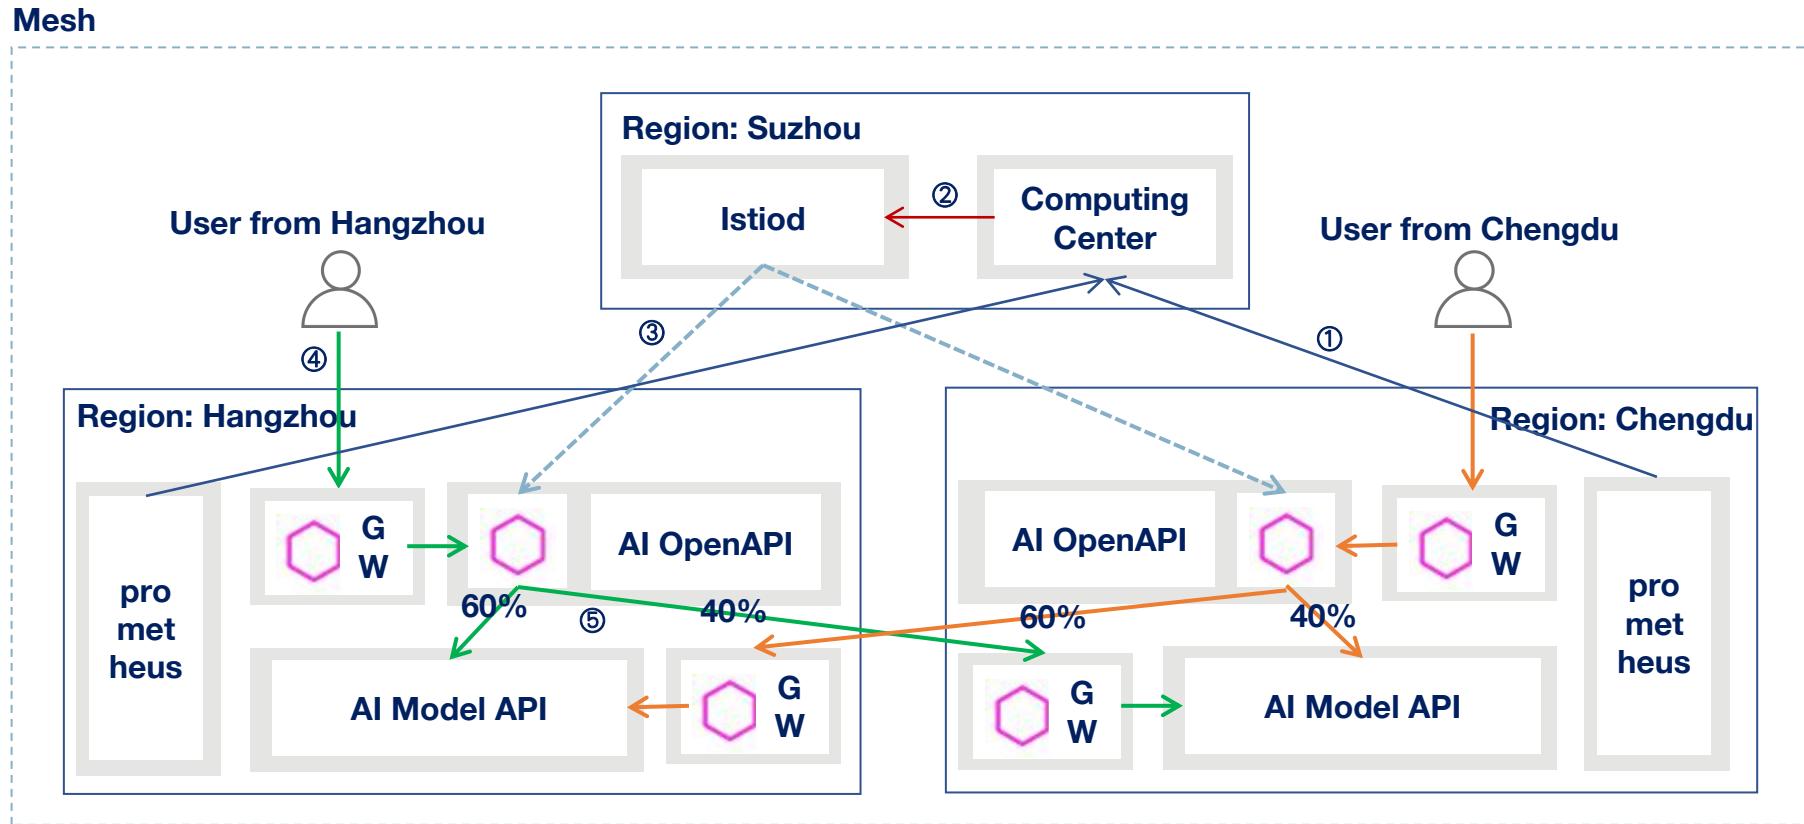
## Computing Network-Native

The computing network presents new challenges to cloud-native. China Mobile Cloud has proposed the concept of ‘Computing Network-Native’ .

# Computing Network-Native

**Demo:** Users access the facial recognition service from nearby locations. The Computing Network-Native dynamically adjusts traffic and allocates computational resources based on resource availability.

**Technology:** Istio Single Mesh Multi-Cluster Multi-Network Model.





KubeCon



CloudNativeCon



OPEN SOURCE SUMMIT

China 2023

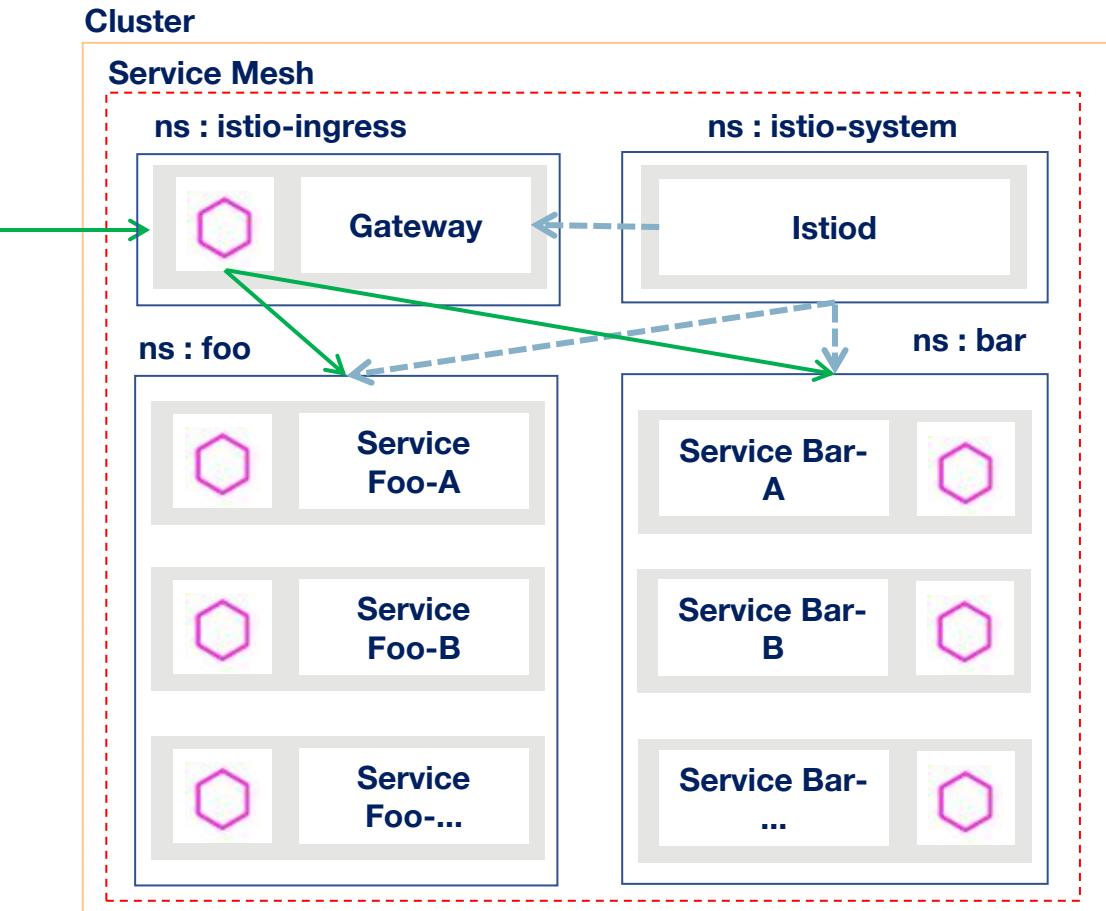
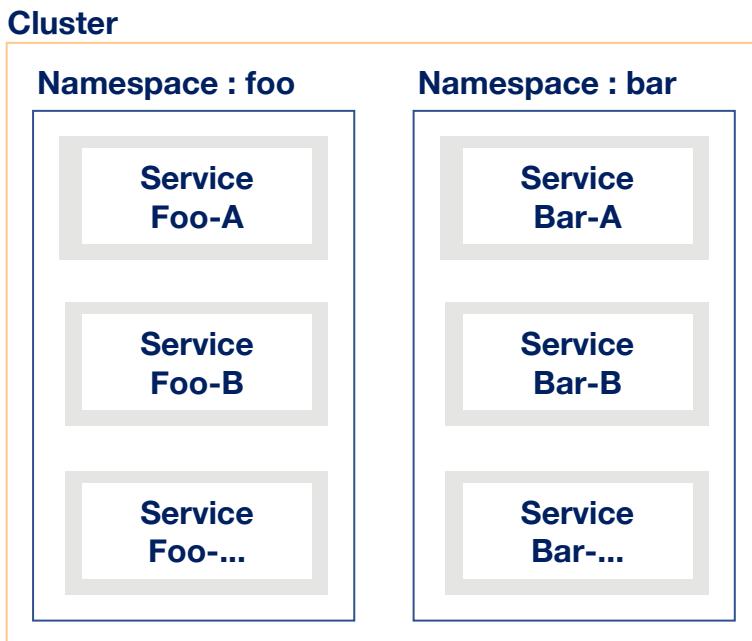
# Technical Exploration

# Multi-Tenant Practice

**Background:** Multiple products are deployed within a single Kubernetes cluster, segmented by namespaces.

**Issue:** If following the Istio classic deployment model, performance and fault isolation issues may arise.

**Requirement:** A multi-tenant solution at the Namespace level for Istio.

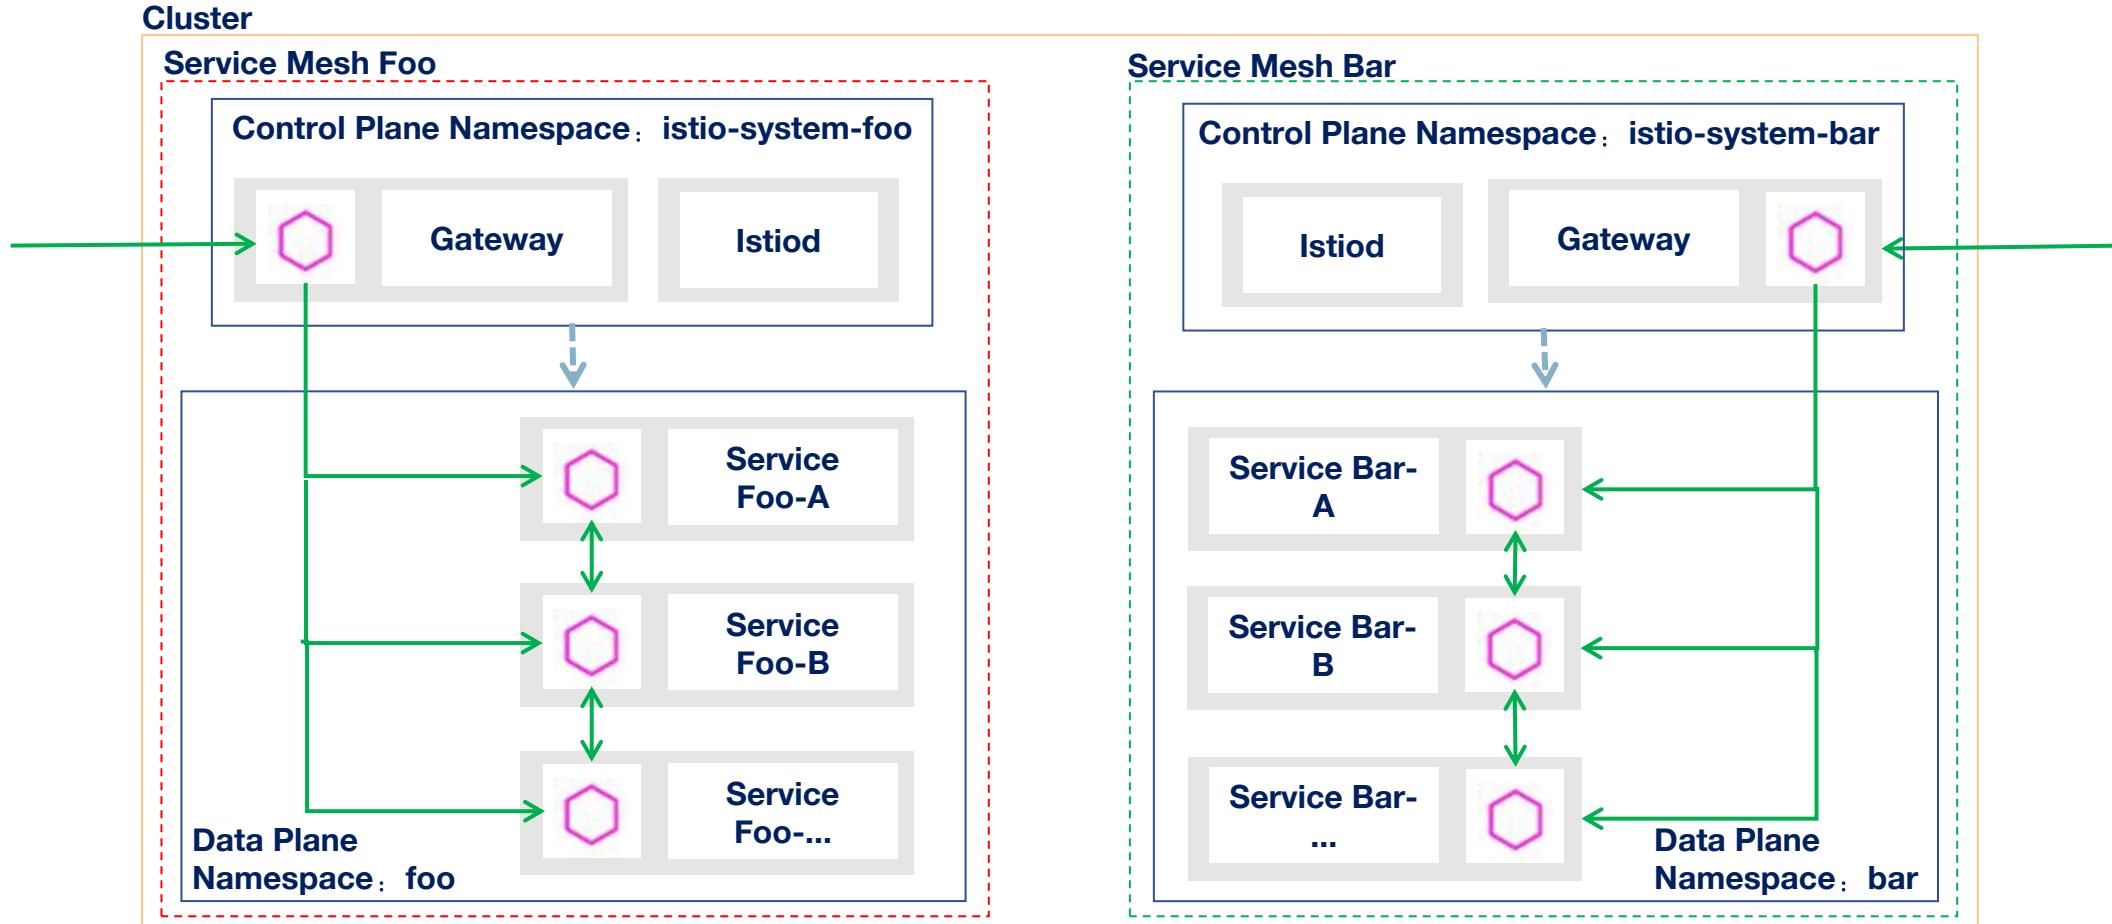


Application Deployment Model

Istio Classic Deployment Model

# Multi-Tenant Practice

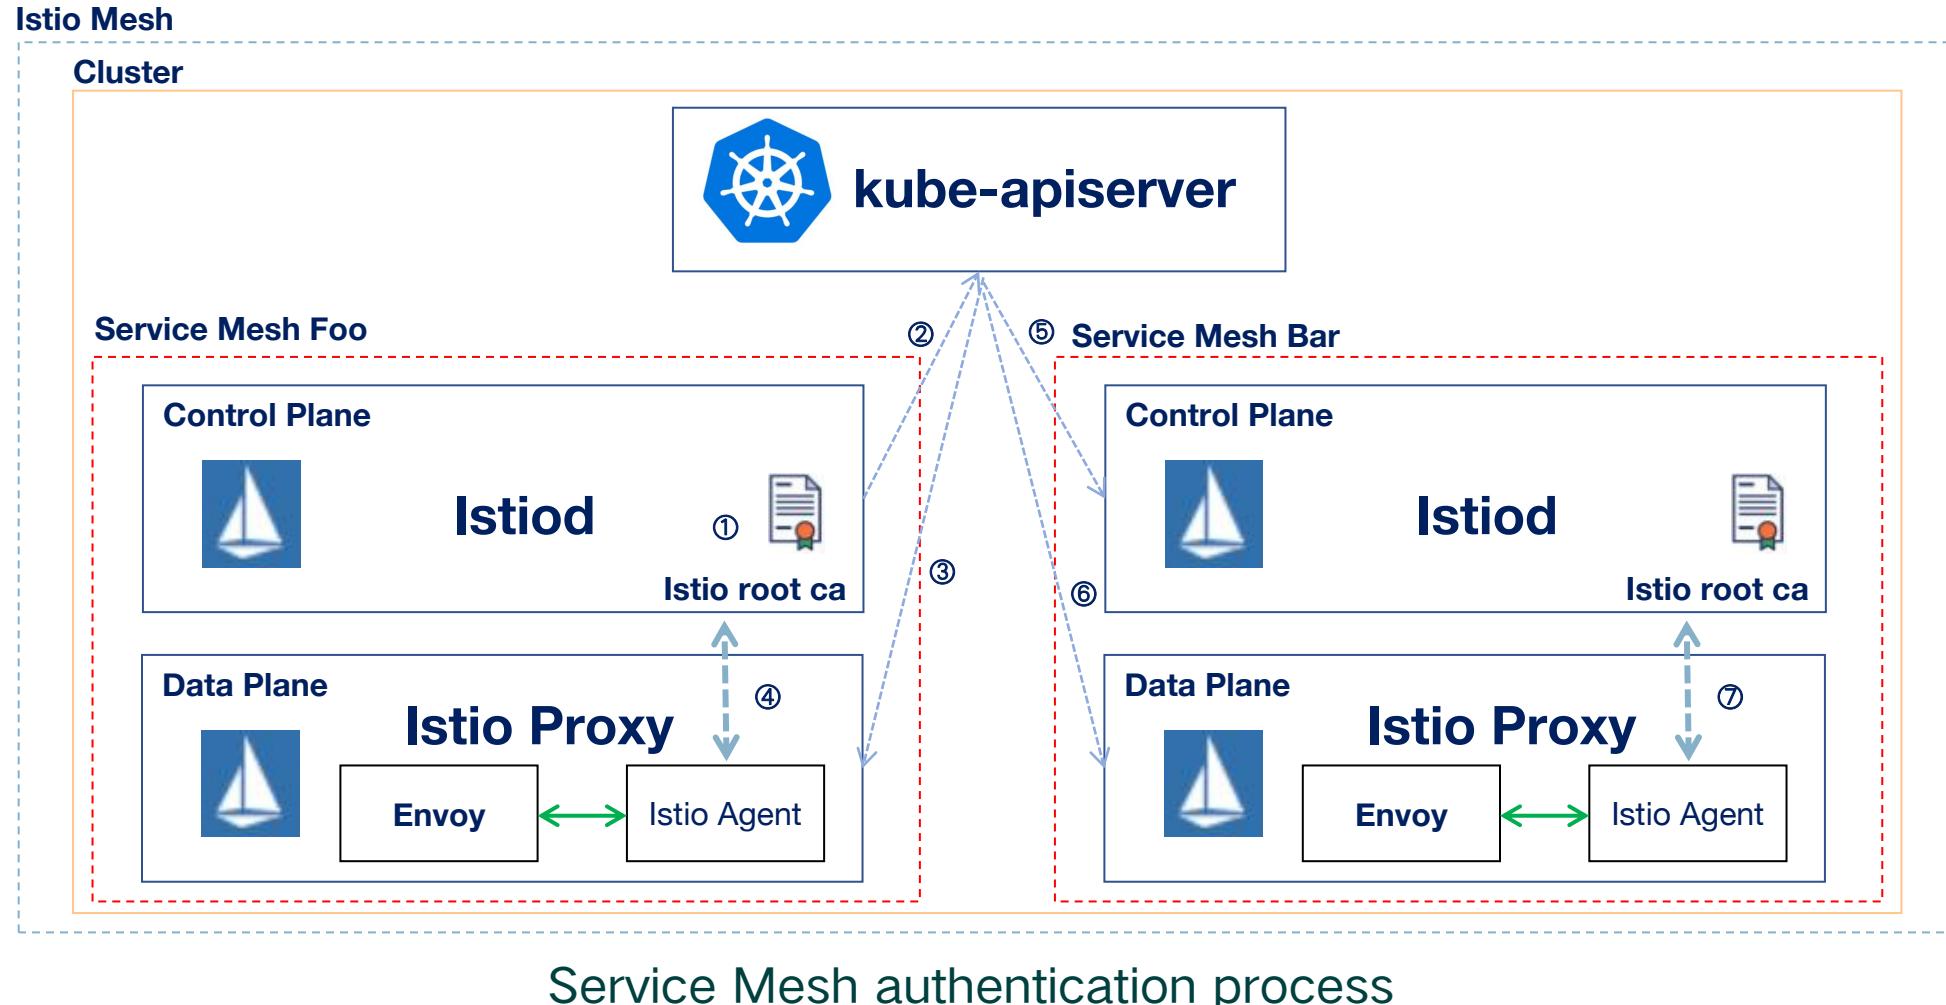
**Ideal Solution:** Tenant isolation for both the control plane and data plane.



Tenant isolation Model for both the control plane and data plane

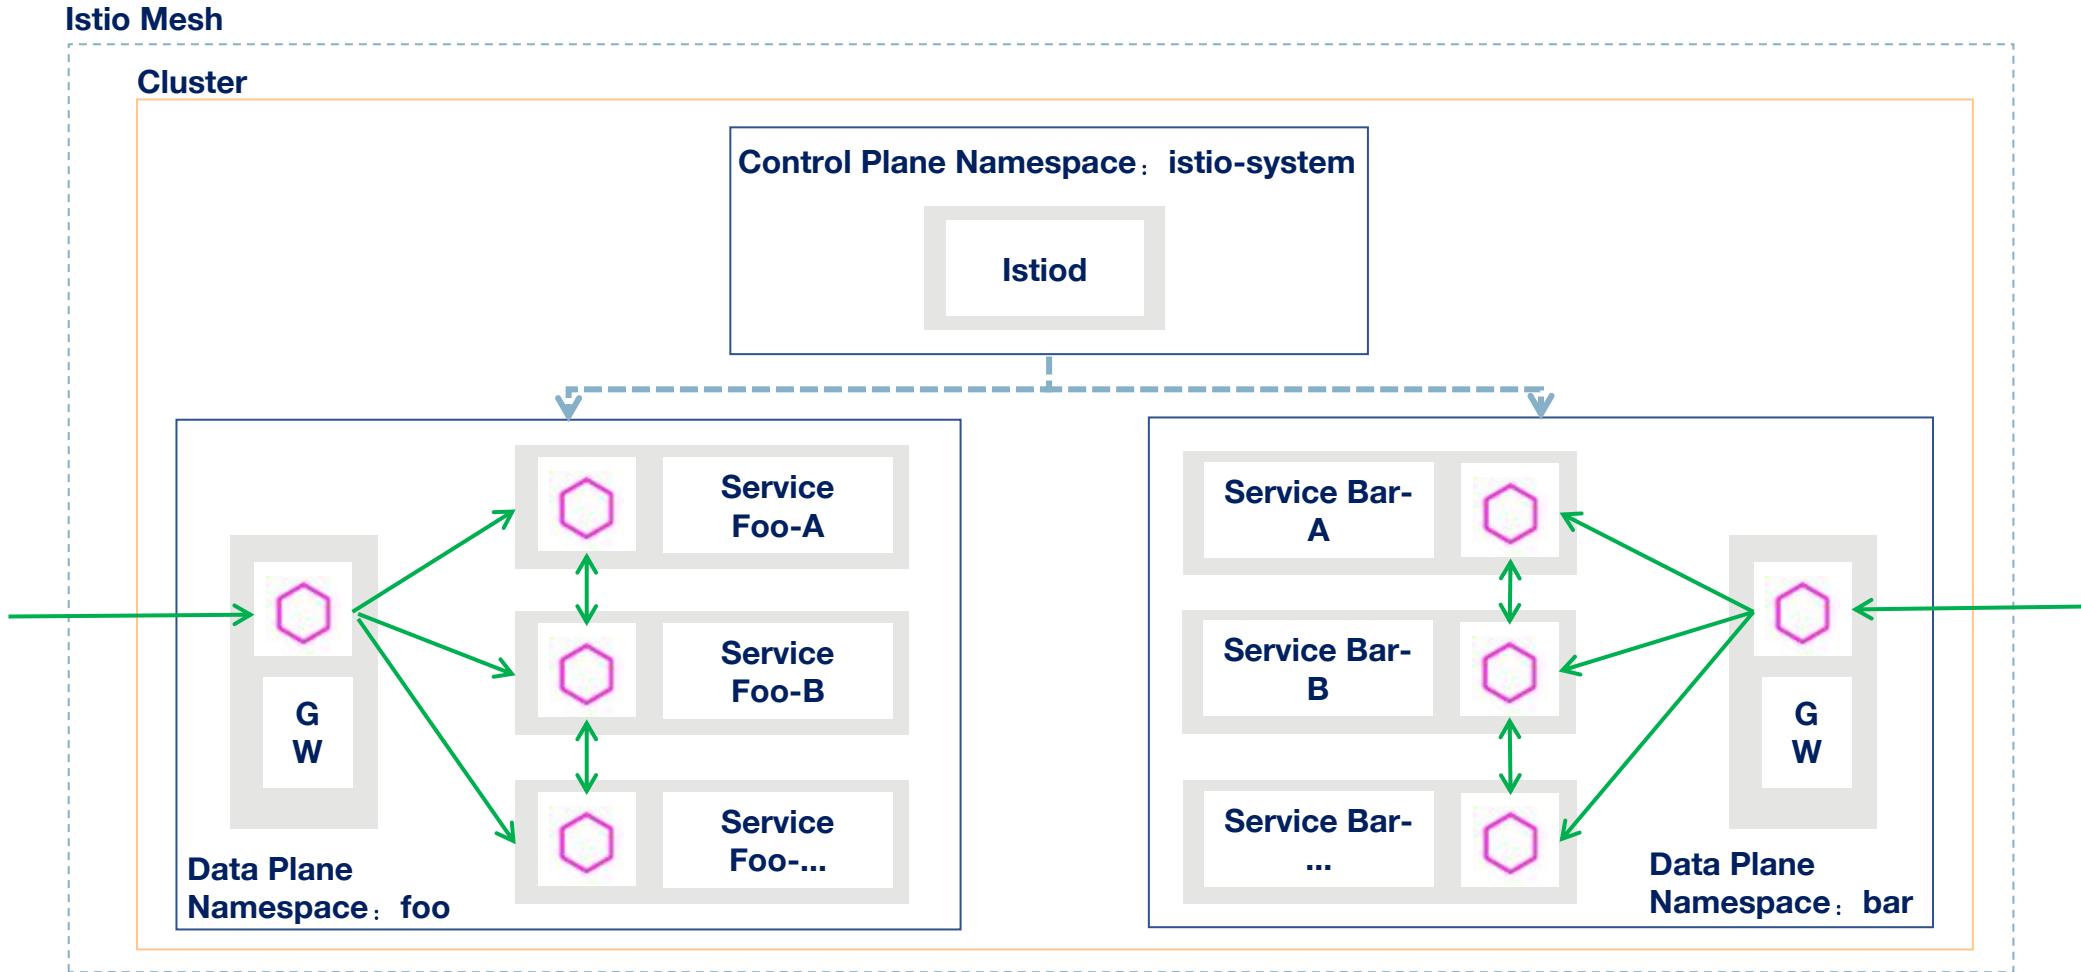
# Multi-Tenant Practice

The solution is feasible, but we decide to abandon it. Because it requires code modifications, leading to high development and maintenance costs.



# Multi-Tenant Practice

**Middle ground Model:** Partial isolation, shared control plane, isolated data plane, no code modification, and no development or maintenance costs. **We have adopted this model.**



# Multi-Tenant Practice

**Key Point:** Istiod's environment variable `PILOT_SCOPE_GATEWAY_TO_NAMESPACE`



```
apiVersion: install.istio.io/v1alpha1
kind: IstioOperator
metadata:
  namespace: istio-system
  name: istio-control-plane
spec:
  namespace: istio-system
  profile: default
  components:
    pilot:
      k8s:
        env:
          - name: PILOT_SCOPE_GATEWAY_TO_NAMESPACE
            value: "true"
    ingressGateways:
      - name: istio-ingressgateway
        enabled: false
```

IstioOperator for control plane



```
apiVersion: install.istio.io/v1alpha1
kind: IstioOperator
metadata:
  namespace: istio-system
  name: foo-ingress-gateway
spec:
  profile: empty
  components:
    ingressGateways:
      - name: istio-ingress-gateway
        namespace: foo
        enabled: true
```

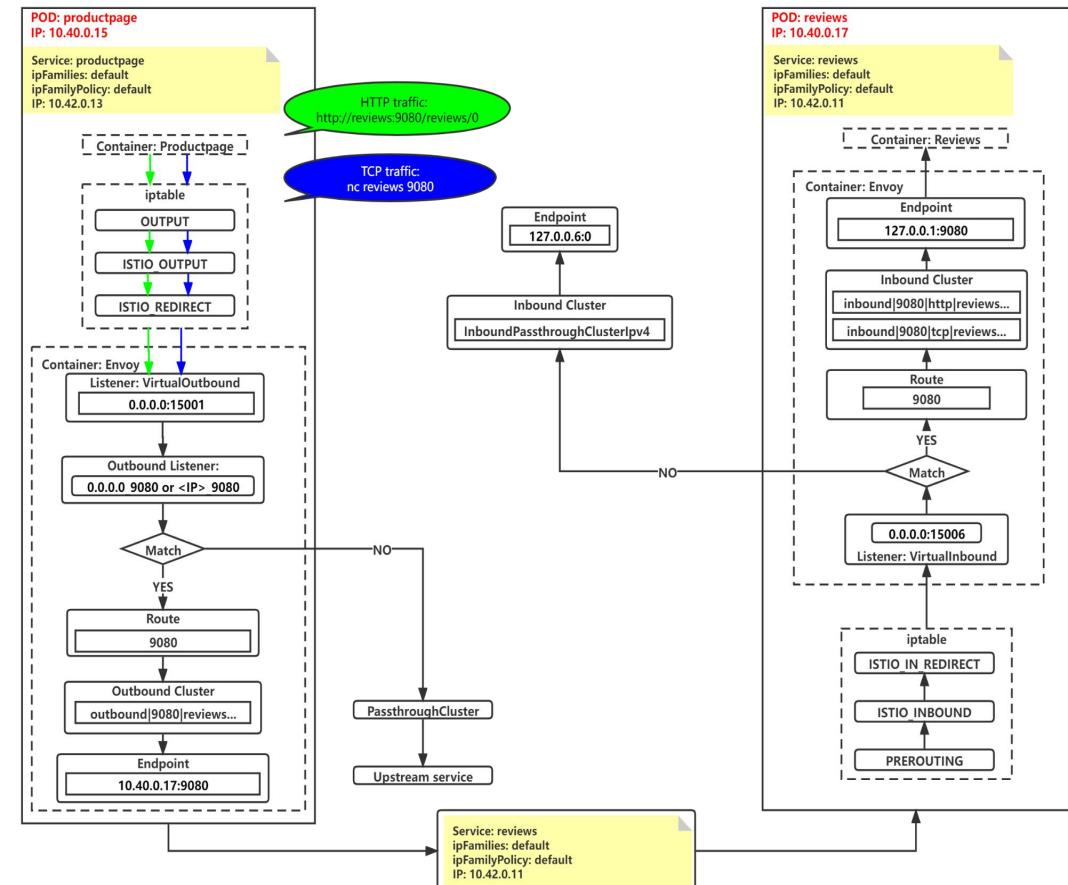
IstioOperator for ingress gateway

# Dual-Stack Practice

With the increasing trend towards IPv6 adoption, In 2022, China Mobile Cloud planned to enable all products to be accessible via both IPv4 and IPv6, implementing a dual-stack approach. However, **Istio** did not support dual-stack network environments in the early part of 2022.

 **Closed** Istio-1.13.2 not support K8S 1.21 dualStack ReadinessProbe failed #38314  
wangyj-tc opened this issue on Apr 10, 2022 · 16 comments

 **howardjohn** commented on Apr 12, 2022  
  
Dual stack is not yet supported, see [#37533](#)  

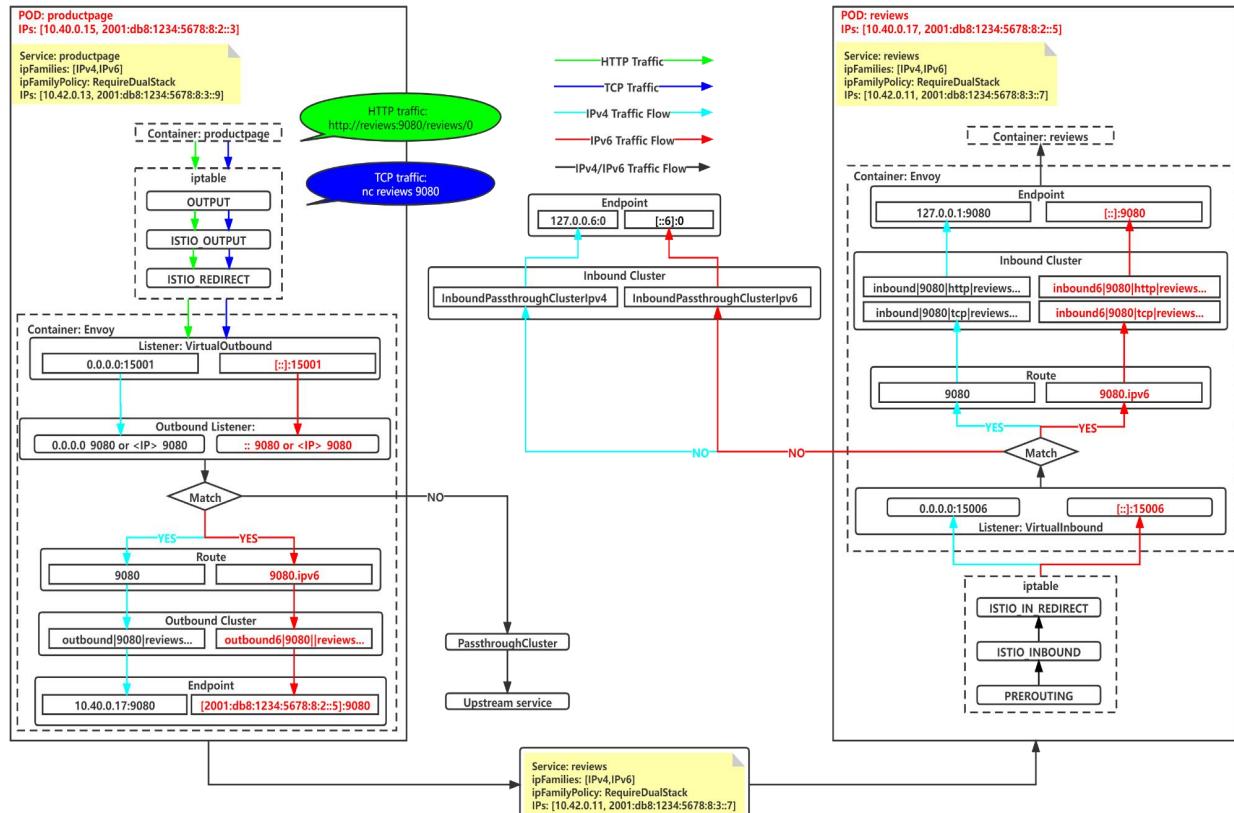



Not supporting dual-stack, only single-stack.

# Dual-Stack Practice

Based on the solution involving the generation of duplicate sets of xDS configurations, We collaborated with Intel to explore and implement dual-stack technology in Istio, contributed exploration findings back to the Istio Community ([Istio experimental-dual-stack branch](#)) .

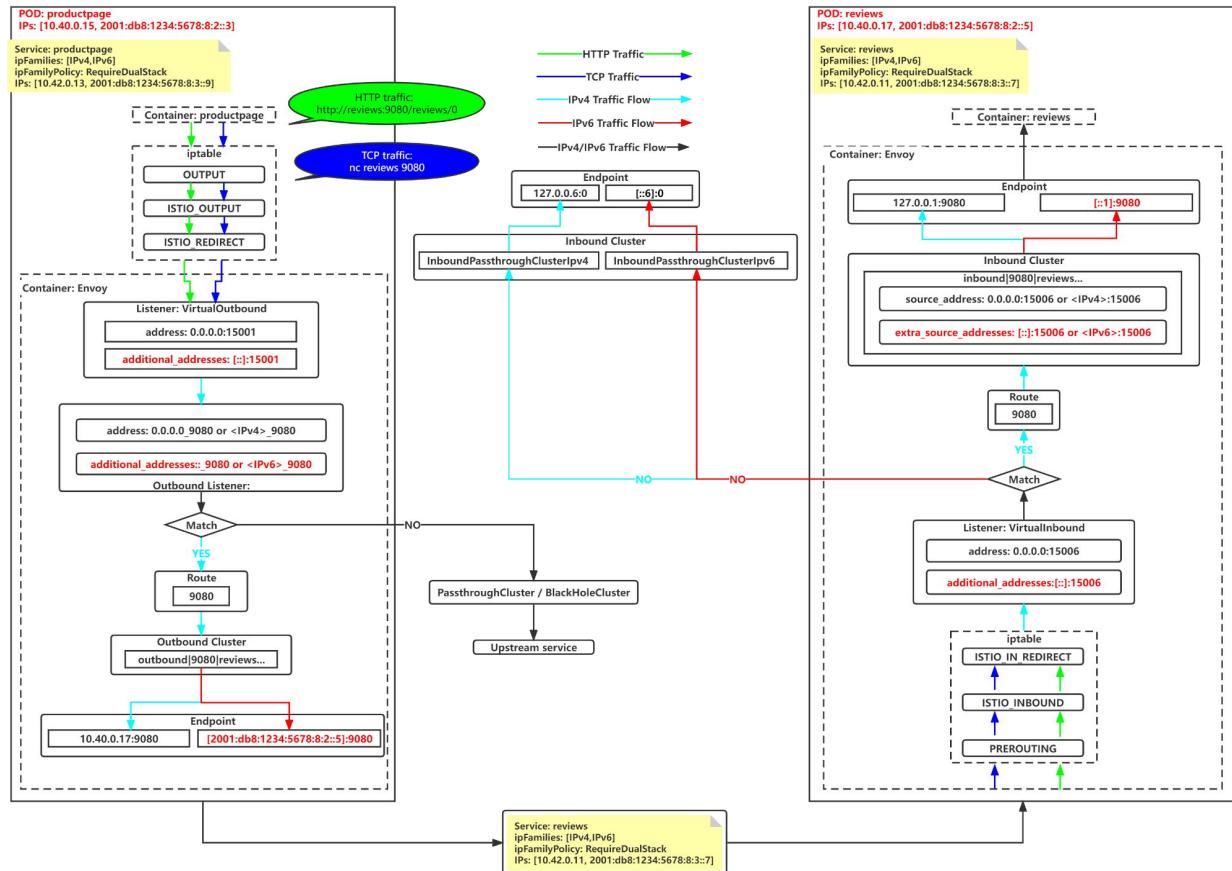
It has been widely adopted within China Mobile Cloud and China Mobile, becoming **the first successful deployment of dual-stack Istio Community solution in China**.



Duplicate sets of xDS configurations solution

# Dual-Stack Practice

Currently, the Istio and Envoy communities are optimizing and enhancing the dual-stack solution. They have achieved a dual-stack solution based on a single set of xDS configurations to reduce resource consumption and improve performance. This enhancement was released in version 1.17. China Mobile Cloud is currently experimenting with this solution.





KubeCon



CloudNativeCon



OPEN SOURCE SUMMIT

China 2023

# Product Development

# Evolution Process

The exploration of China Mobile Cloud's service mesh has undergone an evolutionary process



Usage in component-level

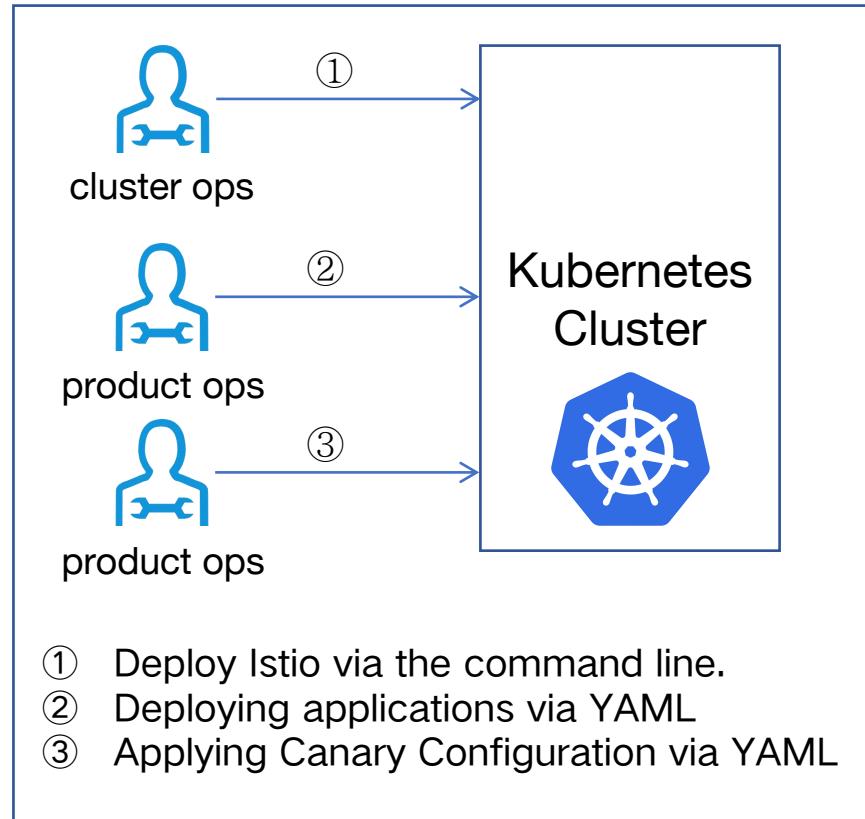


Integration with OPS platforms

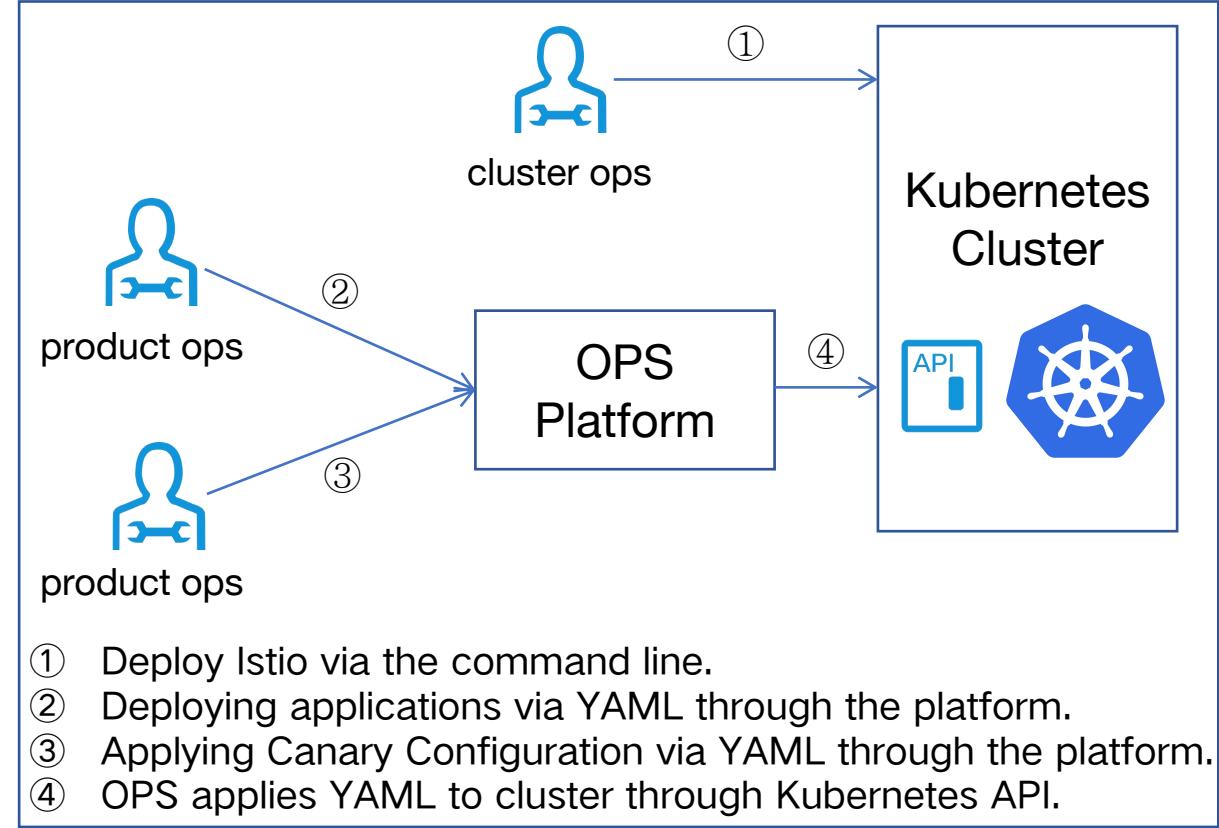


Combining with OAM

# Evolution Process



Usage in component-level



Integration with OPS platforms

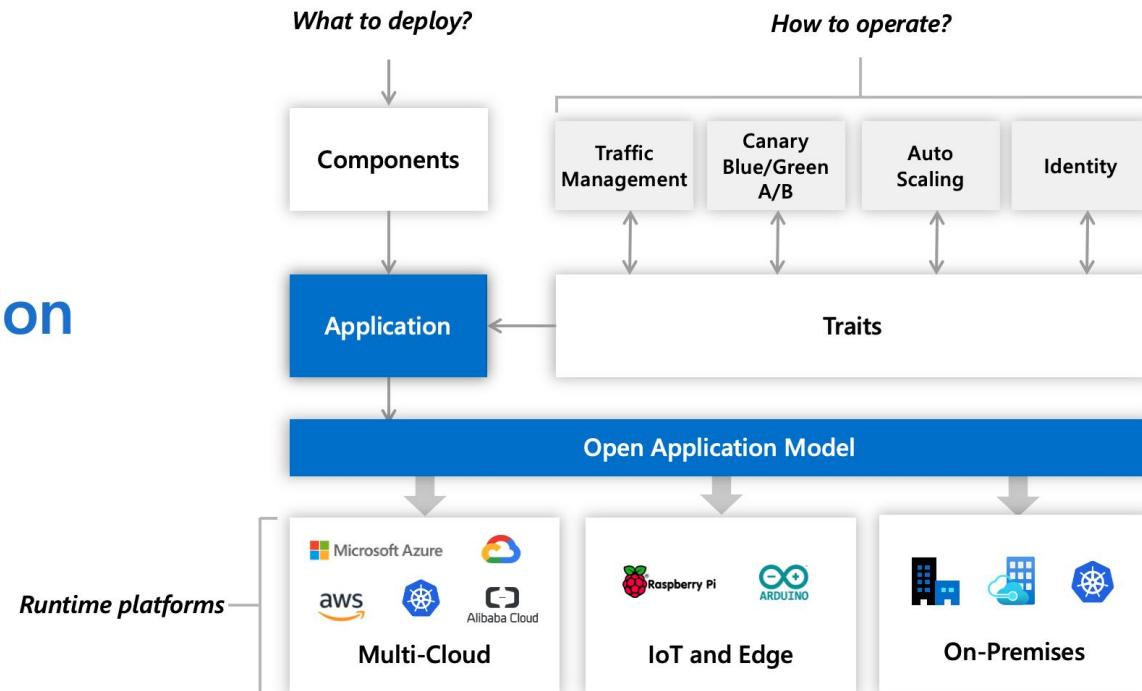
Whether at the component level or through integration with operational platforms, the product teams have been compelled to become Kubernetes experts, Istio experts, and YAML engineers.

We are gradually gaining a clear understanding of the cloud-native application platform, and We need to provide a platform that adheres to the following principles:

- Abstracting the underlying infrastructure.
- User-centric with applications at the core.
- Excellent scalability.

The Open Application Model (OAM) perfectly meets our needs.

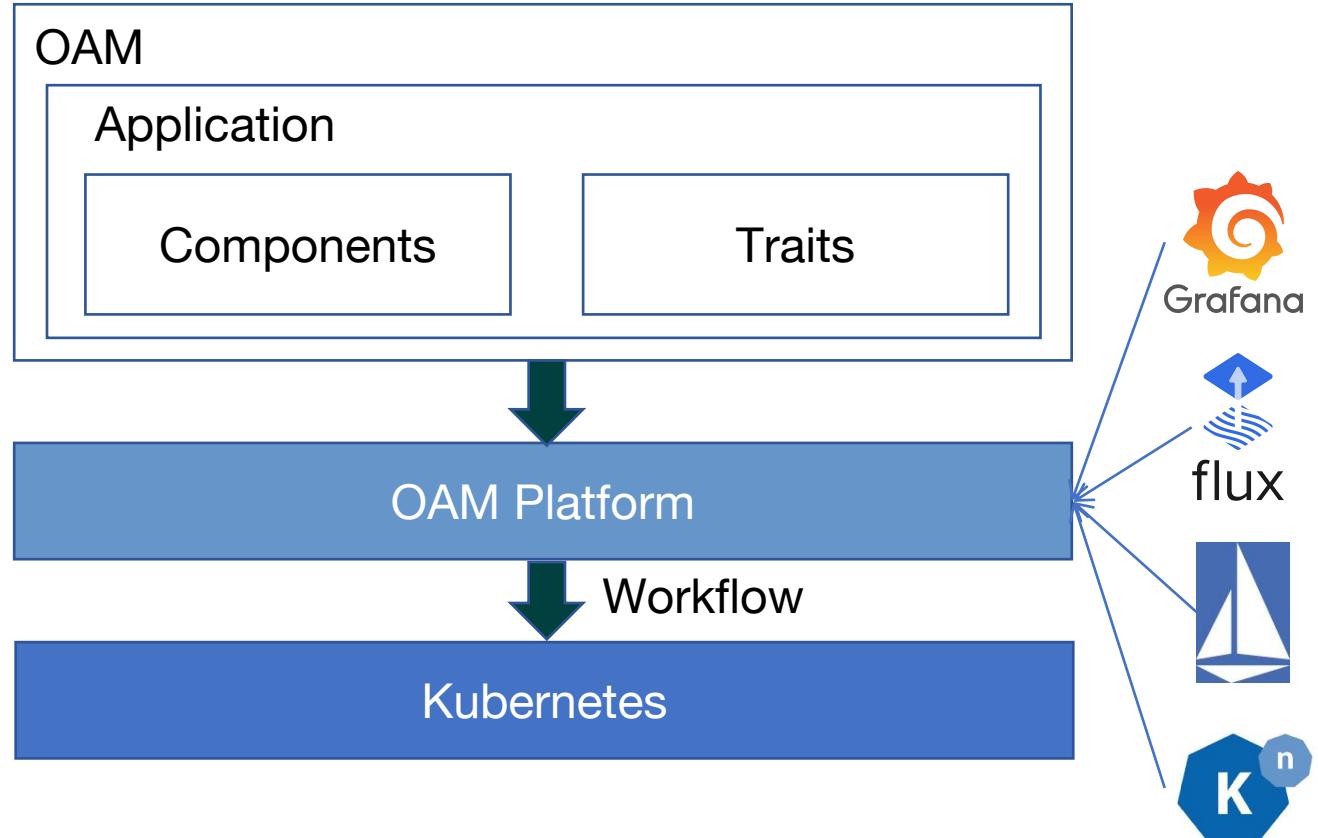
## Open Application Model



# OAM+OAM Platform

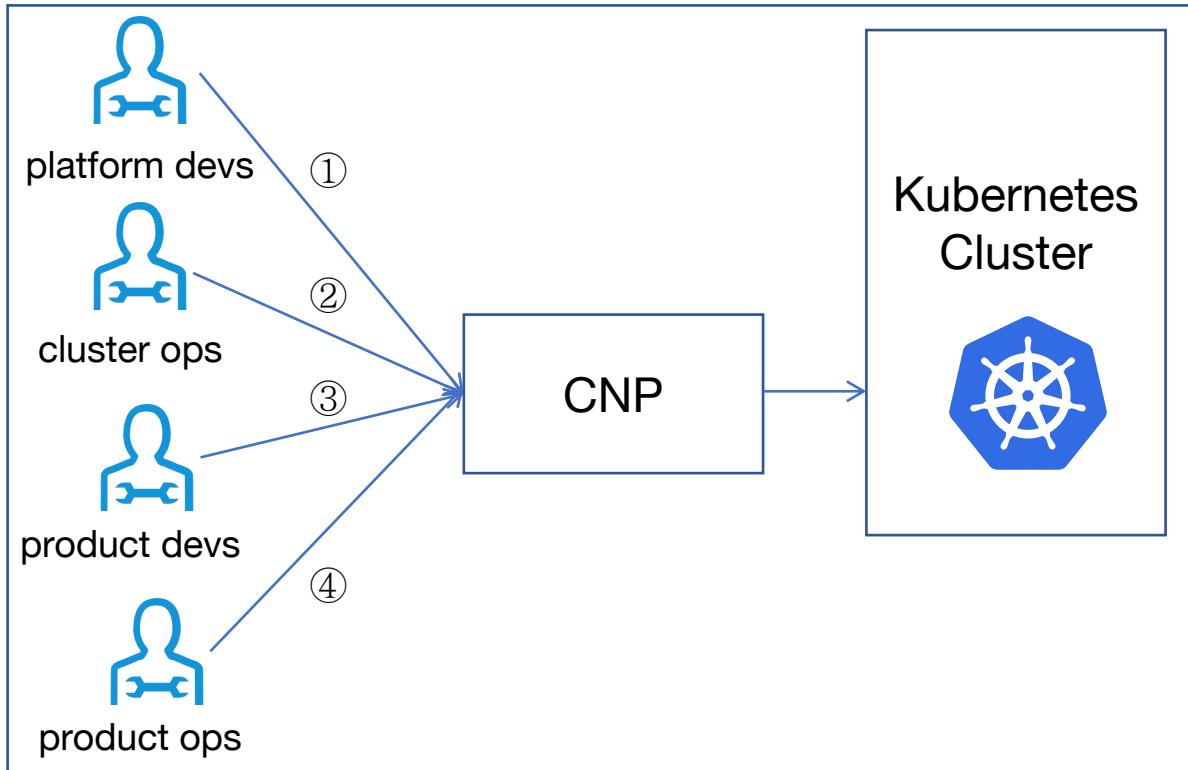
## Basic Concepts:

- **Component**: The fundamental building block for defining an application.
- **Trait**: Operational properties defined on components.
- **Application**: Combining Components and Traits together to represent the complete configuration of an application.
- **OAM Platform**: The platform for implementing the OAM model like KubeVela.
- **Plugin**: The OAM platform introduces some ecosystem components in the form of plugins.
- **Workflow**: Part of the application deployment plan, assisting users in customizing the steps within the application deployment plan.



# Evolution Process

We have explored and built our own cloud-native platform (CNP) based on the OAM and OAM Platform (KubeVela) to serve both internal products within China Mobile Cloud and external customers. We are also actively contributing our exploration findings back to the KubeVela community.



Combining with OAM

- ① Develop Istio plugin, canary release traits, canary release workflows, and integrate them into CNP.
- ② Deploy Istio through the plugin-based approach.
- ③ Create Application Components.
- ④ Creating product-specific canary traits and workflow, deploy canary versions of components and apply grey release configurations through the workflow.

# Technical Implementation

As platform developers, our task is only to develop Istio plugins, canary release traits, canary release workflows, and integrate them into the platform.

```
package main

import "strings"

_base: *"" | string
registry: *"" | string
deployGroup: [] | [...]

if parameter.registry != _|_ {
    registry: parameter.registry
}
if parameter.deployGroup != _|_ {
    deployGroup: parameter.deployGroup
}

if registry != "" && !strings.HasSuffix(registry, "/") {
    _base: registry + "/"
}
if registry == "" || strings.HasSuffix(registry, "/") {
    _base: registry
}

let _deployNamePrefix = "deploy-"
let _overrideNamePrefix = "override-"

// produce deploy policy according to deployGroup parameter
deployPolicy: [ for k, v in deployGroup {
    type: "topology"
    name: _deployNamePrefix + "\$(k)"
    properties: {
        namespace: "istio-system"
        if v.clusters != _|_ {
            clusters: v.clusters
    }
}
```

Istio plugin template

```
# Code generated by KubeVela templates. DO NOT EDIT. Please edit
# iDefinition source cue file: vela-templates/definitions/internal
apiVersion: core.oam.dev/v1beta1
kind: TraitDefinition
metadata:
    annotations:
        definition.oam.dev/description: use istio to manage traffic
    name: canary-traffic
    namespace: vela-system
spec:
    appliesToWorkloads:
        - deployments.apps
    podDisruptive: true
    schematic:
        cue:
            template: |-
                import ("vela/op")
                import "strings"

            parameter: {
                target: {
                    cluster: string
                    namespace: string
                }
                component: string
                versions: {
                    oldVersion: string
                    newVersion: string
                }
                header?: {
                    matchRule: **"exact" | "prefix" | "regex"
                    key: string
                    value: string
                }
                uri?: {
                    matchRule: **"exact" | "prefix" | "regex"
                    path: string
                }
            }
        outputs: virtualService: {
            apiVersion: "networking.istio.io/v1alpha3"
            kind: "VirtualService"
            metadata: {
                name: parameter.originName
                namespace: context.namespace
            }
            spec: {
                hosts: [parameter.originName]
                http: [{route: [
                    {destination: {
                        host: parameter.originName
                        subset: "withoutCanary"
                    }}]}
            }
        }
    }
```

canary release trait

```
apiVersion: core.oam.dev/v1beta1
kind: WorkflowStepDefinition
metadata:
    name: canary-rollout-content
    namespace: vela-system
spec:
    schematic:
        cue:
            template: |-
                import ("vela/op")
                import "strings"

            parameter: {
                target: {
                    cluster: string
                    namespace: string
                }
                component: string
                versions: {
                    oldVersion: string
                    newVersion: string
                }
                header?: {
                    matchRule: **"exact" | "prefix" | "regex"
                    key: string
                    value: string
                }
                uri?: {
                    matchRule: **"exact" | "prefix" | "regex"
                    path: string
                }
            }
        outputs: virtualService: {
            apiVersion: "networking.istio.io/v1alpha3"
            kind: "VirtualService"
            metadata: {
                name: parameter.originName
                namespace: context.namespace
            }
            spec: {
                hosts: [parameter.originName]
                http: [{route: [
                    {destination: {
                        host: parameter.originName
                        subset: "withoutCanary"
                    }}]}
            }
        }
    }
```

canary release workflow based on request contents

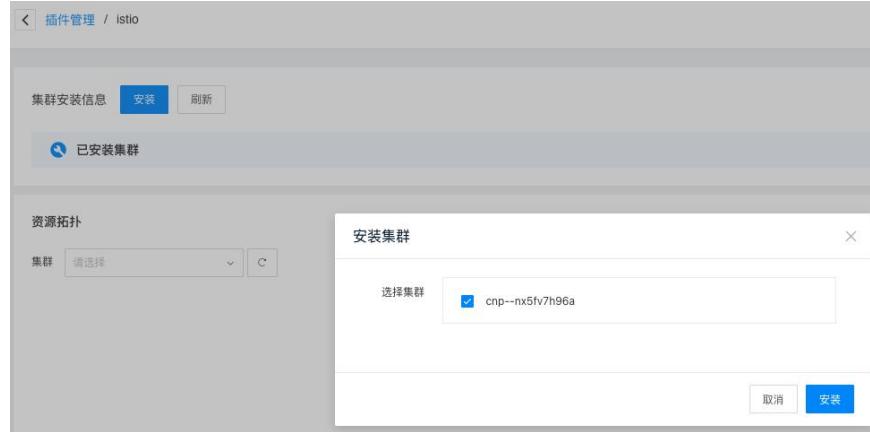
```
apiVersion: core.oam.dev/v1beta1
kind: WorkflowStepDefinition
metadata:
    name: canary-rollout-weight
    namespace: vela-system
spec:
    schematic:
        cue:
            template: |-
                import ("vela/op")
                import "strings"

            parameter: {
                target: {
                    cluster: string
                    namespace: string
                }
                component: string
                traffic: [...{
                    revision: string
                    weight: int
                }]
            }
        outputs: virtualService: {
            apiVersion: "networking.istio.io/v1alpha3"
            kind: "VirtualService"
            metadata: {
                name: parameter.originName
                namespace: context.namespace
            }
            spec: {
                hosts: [parameter.originName]
                http: [{route: [
                    {destination: {
                        host: parameter.originName
                        subset: "withoutCanary"
                    }}]}
            }
        }
    }
```

canary release workflow based on traffic weights

# Product Form

For the product team, there is no need to have in-depth knowledge of Kubernetes and Istio's underlying details. They can easily access service mesh capabilities through the platform.



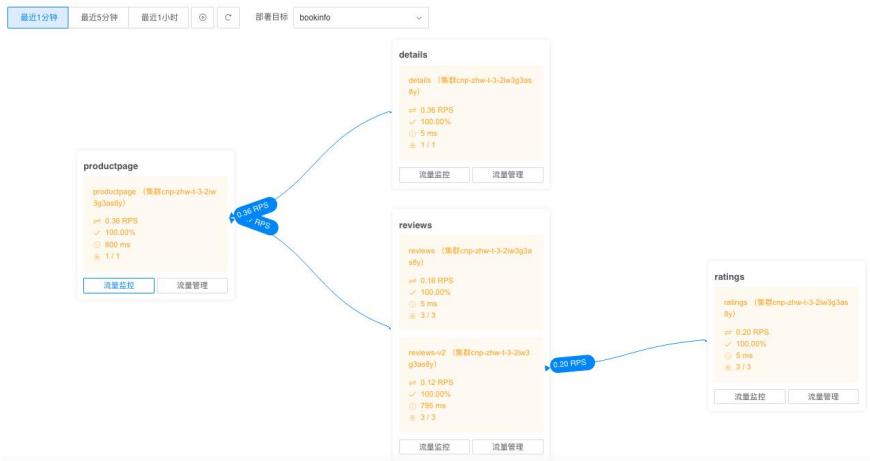
Deploy Istio via plugin

A screenshot of a configuration interface titled "编辑灰度组件" (Edit Gray Component). It includes fields for "名称" (Name: reviews-v2), "描述" (Description: 不超过256个字符), "灰度组件" (Gray Component: reviews), "类型" (Type: webservice), and "依赖" (Dependency: 请选择). Below this is the "应用配置" (Application Configuration) section, which includes a toggle for "服务治理" (Service Governance) set to "开启" (Enabled), a "版本" (Version: v2) field, and a "容器镜像" (Container Image: Docker ... bookinfo/examples-bookinfo-reviews-v2:latest). To the right, another window titled "添加工作流步骤" (Add Workflow Step) is shown, detailing steps for a "灰度发布" (Gray Release) workflow named "canary-review" with target "bookinfo", component "reviews", and policy "完全替换" (Replace All).

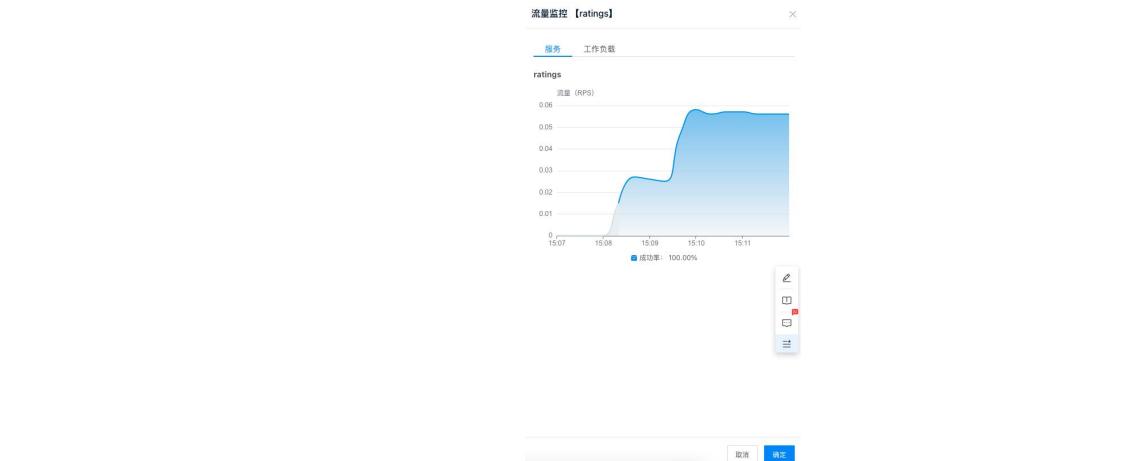
Create canary Component

A screenshot of a configuration interface titled "添加工作流步骤" (Add Workflow Step) for a "灰度发布" (Gray Release) workflow. It shows fields for "工作流类型" (Workflow Type: 灰度发布), "名称" (Name: canary-review), "描述" (Description: 请输入...), "部署目标" (Deployment Target: bookinfo), "组件" (Component: (reviews)), "发布策略" (Release Policy: 全球重发布), and "策略配置" (Policy Configuration: 按流量比例下发). At the bottom, a "完成" (Finish) button is visible.

Define canary workflow



Canary release service topology



Traffic Observability



KubeCon



CloudNativeCon



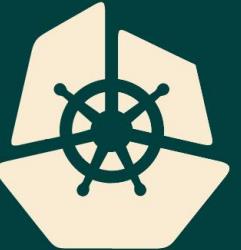
OPEN SOURCE SUMMIT

China 2023

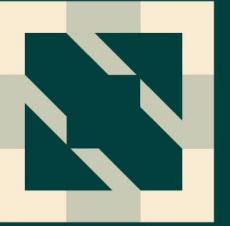
# Futtrue Prospects

# Future Prospects

- Piloting the dual-stack solution based on a single set of xDS configurations.
- Exploring Ambient Mesh
- Serverless Exploration in the Computing Network



KubeCon



CloudNativeCon

S OPEN SOURCE SUMMIT

China 2023